

Parallel Scientific Computation

OpenMP 1

J.-H. Parq

IPCST

Seoul National University

What is OpenMP?

- Open Multi-Processing
 - API(application programming interface) for **shared memory multi-processing**
 - Main function: **Multithreading**
 - Supported languages: Fortran, C, C++
 - Supported platforms: Linux, Windows, macOS, Solaris, AIX, HP-UX,
 - Managed by OpenMP ARB
(Architecture Review Board)




Multithreading

- Thread
 - The smallest unit of a sequence of instructions controlled by a scheduler
- In OpenMP
 - Multi-threads can exist in one process.
 - Threads share resources of the process.
 - Threads can perform concurrent execution.
 - Threads can own their private variables.
 - ❖ Some CPUs provide hyper-threading.

Fork and Join Model

- Branching and merging
 - Fork: A master thread generates slave threads.
 - Join: Slave threads join into the master.



Components

- Compiler directives
- Runtime library routines
- Environment variables

Directives

- Commands meaningful only when the compiler is on the OpenMP mode.
 - The compiler ignores them otherwise.
 - In Fortran: the **comment** form
 - In C or C++: the **pragma** form

Directives

Fortran

- Allowed formats (sentinels)

!\$omp

c\$omp

*\$omp

C / C++

- Allowed formats (sentinels)

#pragma omp

Directives

Fortran

- Example

```
!$omp parallel do  
do I = 1, 60  
    A(I) = A(I)*X(I)  
end do  
!$omp end parallel do
```

C / C++

- Example

```
#pragma omp parallel for  
for (i = 0; i<60; i++) {  
    a(i) = a(i)*x(i)  
}
```


Directives

- Frequently used keywords
 - parallel (+ end)
 - do/for
 - private
 - shared
 - section(s)

Library Routines

- Functions or subroutines used in codes
- Frequently used ones
 - `omp_set_num_threads`
 - `omp_get_num_threads`
 - `omp_get_thread_num`
 - `omp_get_num_procs`
 - `omp_get_max_threads`

External Variables

- Variables used on OS (not in your codes)
- Used by the command setnev(csh/tcsh) or export(sh/bash)
 - Ex.) export OMP_NUM_THREADS=8
- Frequently used variables
 - OMP_NUM_THREADS
 - OMP_SCHEDULE

OMP_THREAD_LIMIT

How to Compile

- Intel

`icc -qopenmp` `icpc -qopenmp`

`ifort -qopenmp`

❖ Alternatively, `-openmp` option can be used although it is deprecated.

- GNU

`gcc -fopenmp` `g++ -fopenmp`

`gfortran -fopenmp`

- PGI

`pgcc -mp`

`pgf90 -mp`

Two-way Compiling

- How to enable either serial or parallel?

- C

- Use of a fake openmp header for serial compiling
 - Or `_OPENMP` macro

Ex.) `#ifdef _OPENMP`
`id = omp_get_thread_num();`
`#endif`

- Fortran

- `!$` or `C$` or `c$` or `*$`

Ex.) `!$ id = omp_get_thread_num()`

Simple Example Program

- **hello.c**

```
#include <stdio.h>
```

```
#include <omp.h>
```

```
main()
```

```
{
```

```
    int id, N;
```

```
    float fraction;
```

```
    N = 4;
```

```
    printf("%d threads set by me\n",N);
```

```
    omp_set_num_threads(N);
```

```
    printf("%d procs\n", omp_get_num_procs());
```

```
    printf("%d max. threads\n", omp_get_max_threads());
```

```
    printf("%d thread now\n\n",omp_get_num_threads());
```

Simple Example Program

```
printf("Fork!\n");  
#pragma omp parallel private(id,fraction) shared(N)  
{  
    id = omp_get_thread_num();  
    printf("Hello, I'm thread %d\n",id);  
    fraction = (float)id/(float)N;  
    printf("%d/%d = %f\n",id,N,fraction);  
}  
printf("Join!\n");  
}
```

Simple Example Program

- **hello.cpp**

```
#include <iostream>
#include <omp.h>
using namespace std;
```

```
main() {
    int id, N;
    float fraction;
```

```
    N = 4;
    cout << N << " threads set by me" << endl;
    omp_set_num_threads(N);
    cout << omp_get_num_procs() << " procs" << endl;
    cout << omp_get_max_threads() << " max. threads" << endl;
    cout << omp_get_num_threads() << " thread now" << endl;
```


Simple Example Program

```
cout << "Fork!" << endl;
#pragma omp parallel private(id,fraction) shared(N)
{
    id = omp_get_thread_num();
    fraction = (float)id/(float)N;

    #pragma omp critical (printing)
    {
        cout << "Hello, I'm thread " << id << endl;
        cout << id << "/" << N << " = " << fraction << endl;
    }
}
cout << "Join!" << endl;
}
```

Simple Example Program

- **hello.f**

```
PROGRAM HELLO
```

```
IMPLICIT NONE
```

```
INCLUDE "omp_lib.h"
```

```
INTEGER ID, N
```

```
REAL FRACTION
```

```
N = 4
```

```
PRINT '(I2,A)',N," threads set by me"
```

```
CALL OMP_SET_NUM_THREADS(N)
```

```
PRINT '(I2,A)',OMP_GET_NUM_PROCS (), " procs"
```

```
PRINT '(I2,A)',OMP_GET_MAX_THREADS()," max. threads"
```

```
PRINT *
```

```
PRINT '(I2,A)', OMP_GET_NUM_THREADS() , " thread now"
```

Simple Example Program

```
PRINT '(A)', "Fork!"
```

```
C$OMP PARALLEL
```

```
C$OMP& PRIVATE(ID,FRACTION) SHARED(N)
```

```
ID = OMP_GET_THREAD_NUM ()
```

```
PRINT '(A,I2)', "Hello, I'm thread", ID
```

```
FRACTION = REAL(ID)/REAL(N)
```

```
PRINT '(I2,A,I1,A,F5.2)', ID, "/", N, " = ", FRACTION
```

```
C$OMP END PARALLEL
```

```
PRINT '(A)', "Join!"
```

```
STOP
```

```
END
```

Simple Example Program

- **hello.f90**

```
PROGRAM HELLO
```

```
  USE OMP_LIB
```

```
  IMPLICIT NONE
```

```
  INTEGER ID, N
```

```
  REAL FRACTION
```

```
  N = 4
```

```
  PRINT '(I2,A)',N," threads set by me"
```

```
  CALL OMP_SET_NUM_THREADS(N)
```

```
  PRINT '(I2,A)',OMP_GET_NUM_PROCS (), " procs"
```

```
  PRINT '(I2,A)',OMP_GET_MAX_THREADS()," max. threads"
```

```
  PRINT *
```

```
  PRINT '(I2,A)', OMP_GET_NUM_THREADS() , " thread now"
```

Simple Example Program

```
PRINT '(A)', "Fork!"
```

```
!$OMP PARALLEL &
```

```
!$OMP PRIVATE(ID,FRACTION) SHARED(N)
```

```
  ID = OMP_GET_THREAD_NUM ()
```

```
  PRINT '(A,I2)', "Hello, I'm thread", ID
```

```
  FRACTION = REAL(ID)/REAL(N)
```

```
  PRINT '(I2,A,I1,A,F5.2)', ID, "/", N, " = ", FRACTION
```

```
!$OMP END PARALLEL
```

```
PRINT '(A)', "Join!"
```

```
STOP
```

```
END PROGRAM HELLO
```

Simple Example Program

- **Execution**

g++ hello.cpp -fopenmp -o a.out
./a.out

4 threads set by me

4 procs

4 max. threads

1 thread now

Fork!

Hello, I'm thread 0

$0/4 = 0$

Hello, I'm thread 1

$1/4 = 0.25$

Hello, I'm thread 2

$2/4 = 0.5$

Hello, I'm thread 3

$3/4 = 0.75$

Join!

Parallel Directive

- Block structure
 - Fork at the start and join at the end
 - Synchronization at the end: implicit barrier

- Fortran (f90)

```
!$omp parallel [clauses]
```

```
.....
```

```
!$omp end parallel
```

- C/C++

```
#pragma omp parallel [clauses]
```

```
{
```

```
.....
```

```
}
```

Reduction Clause

- Results of all threads reduce to the master

- Ex.)

```
!$omp parallel reduction(+:b) private(i) shared(a)
```

```
    i = omp_get_thread_num()
```

```
    b = b + a(i)*a(i)
```

```
!$omp end parallel
```

- About the operation

- Addition(+:...), multiplication(*:...),
- 0(+) or 1(*) for the initial value of the variable which all the results reduce into

Private/Shared Clauses

- Private
 - Every thread has an **individual** variable of the same name.
 - **Different memory location**
- Shared
 - **All** thread share the variable.
 - **Same memory location**
 - *Race condition* can form.
 - Many shared variables reduce performance speed.

Race Condition

- Simultaneous access
 - Hazardous if the result depends on the sequence of processes

- Example

```
!$omp parallel private(i) shared(a,b)
```

```
    i = omp_get_thread_num()
```

```
    b = b + a(i)*a(i)
```

```
!$omp end parallel
```

Reading and writing 'b' are problematic!

- A race condition can be solved by adding variables or using directives for synchronization.

Critical Directive

- One by one (to avoid conflicts)
 - Only one thread at a time (not in number order but randomly) executes the commands under 'critical' directive.

- Fortran (f90)

```
!$omp critical [(name)]
```

```
.....
```

```
!$omp end critical [(name)]
```

- C/C++

```
#pragma omp critical [(name)]
```

```
.....
```

```
or {
```

```
.....
```

```
}
```

Critical Directive

- Notes
 - Use only if indispensable
 - Total computational time increases.
 - But necessary if race conditions emerge.
 - Naming is important.
 - All 'critical' blocks with no name are regarded as the same block.
 - If one thread enter a 'critical' block, the other threads **cannot** enter **any** 'critical' blocks until the thread comes out of the block.

Atomic Directive

- Similar to 'critical' but only for atomic operations.
- No block structure.
- Faster than 'critical'.

- Fortran (f90)

```
!$omp atomic
```

```
.....
```

- C/C++

```
#pragma omp atomic
```

```
.....
```

Atomic Directive

- Atomic operations in OpenMP

$x++$ $x--$ (for C/C++)

$x \mathrel{\mathbb{b}} = \textit{expre}$ (for C/C++)

$x = x \mathrel{\mathbb{b}} \textit{expre}$

$x = \textit{expre} \mathrel{\mathbb{b}} x$

- \mathbb{b} : binary operator (+, -, *, /,). Not overloaded.
- *expre*: expression without x
- ❖ x must be a scalar variable.

Other Synchronization Directives

- Barrier
 - Threads wait until all threads reach the point.
 - Some directives already include implicit barriers and do not need barrier directive.
- Master
 - Only the master executes the command block while the other threads skip it. (No barrier)

References

- C. Evangelinos,
Parallel Programming for Multicore
Machines Using OpenMP and MPI
- B. Barney,
OpenMP
- Wikipedia