

# Parallel Scientific Computation

## OpenMP 3


J.-H. Parq

IPCST

Seoul National University

# Numerical Integration

- Trapezoid rule



$$\frac{\Delta x}{2} [f_0 + 2f_1 + 2f_2 + \cdots + 2f_{N-1} + f_N]$$

– Ex.)

$$\int_0^1 \frac{4}{1+x^2} dx = \pi$$

# Numerical Integration

## Fortran

```
dx = 1d0/dble(N)
S = S + 4d0/2d0
!$omp parallel private(x)
!$omp do reduction(+:S)
    do I = 1, N-1
        x = dble(I)*dx
        S = S + 4d0/(1d0 + x*x)
    end do
!$omp end do
!$omp end parallel
S = S + 2d0/2d0
ANSWER = S*dx
```

## C

```
dx = 1.0/(double)N
s = 4.0/2.0;
#pragma omp parallel private(i,x)
{
    #pragma omp for reduction(+:s)
    for (i=1; i<N; i++) {
        x = (double)i*dx;
        s += 4.0/(1.0 + x*x);
    }
}
s += 2.0/2.0;
answer = s*dx;
```

# Loop Dependency

- Loop-independent dependence
  - Dependence exists within the same iteration
  - No hindrance in parallelization
- Loop-carried dependence
  - Dependence between different iterations
  - Ex.)  
    for (i=1; i<N; i++) a[i] += a[i-1];
  - Difficult to parallelize

# Loop Dependency

- Removing loop-carried dependence

- Several cases

1. Loop skewing

- Serial example

```
/*Arrays have already been assigned.*/
```

```
for (i=1; i<N; i++) {  
    b[i] += a[i-1];  
    a[i] += c[i];  
}
```

- Parallel version

```
b[1] += a[0];  
#pragma omp for private(i)  
for (i=1; i<N-1; i++) {  
    a[i] += c[i];  
    b[i+1] += a[i];  
}  
a[N-1] += c[N-1];
```

# Loop Dependency

- Removing loop-carried dependence

## 2. Creating a clone array

- Serial example

```
/*Array a has already been assigned.*/
```

```
for (i=0; i<N-1; i++) {  
    tmp = (b[i]+c[i])/2.0;  
    a[i] = a[i+1]+tmp;  
}
```

- Parallel version

```
#pragma omp for private(i)  
for (i=0; i<N-1; i++) ta[i] = a[i+1];
```

```
#pragma omp for private(i,tmp)  
for (i=0; i<N-1; i++) {  
    tmp = (b[i]+c[i])/2.0;  
    a[i] = ta[i]+tmp;  
}
```

# Loop Dependency

- Removing loop-carried dependence

## 3. Loop fission

- Serial example

```
for (i=1; i<N; i++) {  
    a[i] = a[i-1]+a[i];  
    tot += (b[i]+c[i])/2.0;  
}
```

- Parallel version

```
for (i=1; i<N; i++) {  
    a[i] = a[i-1]+a[i];  
}
```

```
#pragma omp for reduction(+:tot)  
private(i)  
for (i=1; i<N; i++) {  
    tot += (b[i]+c[i])/2.0;  
}
```

# Loop Dependency

- Removing loop-carried dependence
  - 4. Using reduction clause
    - Already showed in the integration example
  - 5. Summation  $\rightarrow$  direct calculation
    - Ex.)  $\text{sum}(0 \text{ to } n) \rightarrow n(n+1)/2$



# Time Check

- You must check execution time to see how parallelization performance change
  1. Command 'time' in Unix or Linux
    - Total elapsed time (*i.e.*, master thread time)
    - Usage) time a.exe
  2. Time library in C
    - Some Fortran versions also provide time library.
    - It is safe to use it only for the master thread
  3. OpenMP routines
    - `omp_get_wtime()`: wall clock time on a thread

# Enhancing Efficiency

- Efficiency of *OpenMP* parallelism depends on use of **cache**.

## 1. Number of threads

✓ Just recommendation

- **Even:** Use of L2 cache suppresses overhead increase.
- **(# of cores per CPU) × n:** Use of L3 cache suppresses overhead increase.

## 2. Algorithm improvement

- **Rearranging loops:** Loop fusion, loop tiling, .....
- **Reducing other overhead factors**
- **Balancing loads:** schedule clause, .....

# Enhancing Efficiency

- Multi-dimensional array
  - In C, a more precedent index has a larger unit for memory address.

- Better access

```
#pragma omp for
for (i=0; i<N; i++) {
    sum = 0.0;
    for (k=0; k<M; k++)
        sum += A[i][k]*x[k];
    y[i] = sum;
}
```

- Worse access

```
#pragma omp for
for (i=0; i<N; i++) {
    sum = 0.0;
    for (k=0; k<M; k++)
        sum += x[k]*A[k][i];
    y[i] = sum;
}
```

# Enhancing Efficiency

- Multi-dimensional array
  - In Fortran, usually, a more precedent index has a smaller unit for memory address.

- Better access

```
!$omp do
  do i = 1, N
    sum = 0d0
    do k = 1, M
      sum = sum + x(k)*A(k,i)
    end do
    y(i) = sum
  end do
!$omp end do
```

- Worse access

```
!$omp do
  do i = 1, N
    sum = 0d0
    do k = 1, M
      sum = sum + A(i,k)*x(k)
    end do
    y(i) = sum
  end do
!$omp end do
```

# Enhancing Efficiency

- Multi-dimensional array
  - However, memory access depends on the system (OS, compiler, .....).
  - You should test a large array before arranging the order of access.

# Enhancing Efficiency

- Loop fusion
  - Overheads should be reduced.
  - An array in cache should be reused if possible.

- Before

```
#pragma omp for
for (i=0; i<N; i++) a[i] += b[i];
#pragma omp for
for (k=0; k<M; k++) {
    y[k] *= y[k];
    z[k] = a[k]*k;
}
```

- After

```
#pragma omp for
for (k=0; k<M; k++) {
    y[k] *= y[k];
    if (k<N) a[k] += b[k];
    z[k] = a[k]*k;
}
```

# Enhancing Efficiency

- Collapse clause
  - Fusing nested loops
  - Reduced overheads (and easier to balance)

- Fortran

```
!$omp do collapse(2)
  do J = 1, M
    do I = 1, N
      A(I,J) = A(I,J)*2d0
    end do
  end do
!$omp end do
```

- C/C++

```
#pragma omp for collapse(2)
for (i=0; i<N; i++) {
  for (j=0; j<M; j++)
    a[i][j] *= 2.0;
}
```

# Enhancing Efficiency

- Loop fission
  - An array in cache should be reused if possible.

- Before

```
#pragma omp for
for (i=0; i<N; i++) {
    y[i] = exp(-(double)i);
    for (k=0; k<M; k++) {
        a[k][i] += b[k][i] + c[i];
    }
}
```

- After

```
#pragma omp for
for (i=0; i<N; i++)
    y[i] = exp(-(double)i);
#pragma omp for
for (k=0; k<M; k++)
    for (i=0; i<N; i++)
        a[k][i] += b[k][i] + c[i];
```



# Enhancing Efficiency

- Loop tiling
  - 2-D array  $\rightarrow$  tiles of data chunks fit for cache
  - Rearranging loops for tiles
  - Ex.) if cache size =  $Csize\_x^2 \times 2$ ,  
for (j1=0; j1<N; j1+=Csize\_x)  
  for (i=0; i<N; i++)  
    for (j2=0; j2<N-j1 && j2<Csize\_x; j2++)  
      b[i][j1+j2] = a[j1+j2][i] ;

# Enhancing Efficiency

- Relocating parallel regions

1. Merging

- Before

```
#pragma omp parallel for  
for (i=0; i<N; i++) {  
    .....  
}  
#pragma omp parallel for  
for (k=0; k<M; k++) {  
    .....  
}
```

- After

```
#pragma omp parallel  
{  
    #pragma omp for nowait  
    for (i=0; i<N; i++) {  
        .....  
    }  
    #pragma omp for  
    for (k=0; k<M; k++) {  
        .....  
    }  
}
```

# Enhancing Efficiency

- Relocating parallel regions

## 2. Expanding

- Before

```
for (i=0; i<N; i++) {  
  #pragma omp parallel for  
  for (k=0; k<M; k++) {  
    .....  
  }  
}
```

- After

```
#pragma omp parallel  
{  
  for (i=0; i<N; i++) {  
    #pragma omp for nowait  
    for (k=0; k<M; k++) {  
      .....  
    }  
  }  
}
```


# Enhancing Efficiency

- If clause
  - It is difficult to speed up by parallelizing loops of small size.
  - The if clause checks the loop size and determines if the loop will be parallelized.
  - Ex.)

```
#pragma omp parallel for if (M>threshold)
  for (k=0; k<M; k++) {
    .....
  }
```

# Task Directive

- A task directive just pushes a task into a queue



# Task Directive Example

```
op = 1; i = 0;
#pragma omp parallel shared(a, coord, op, target)
{
    #pragma omp master
    {
        while (op && i<50) {
            #pragma omp task firstprivate(i) private(j)
            {
                for (j=0; j<10000; j++) {
                    if (a[i][j]==target) {
                        coord[0] = i; coord[1] = j; op = 0; } }
            }
            i++;
        }
    }
}
```

# Other Directives

- Single
  - Equivalent to 'sections' with one 'section'
- Workshare
  - Only for Fortran
  - To cover f90 or higher (ex.: forall, where, .....)
- Threadprivate
  - Privatization of global variables or data
  - The 'copyin' clause initializes the variables.

# Other Directives

- Flush
  - Identifying a synchronization point for memory access consistency
  - ‘Flush’ is virtually included in (both explicit and implicit) barriers and ‘critical’.
- Taskwait
  - Barrier for execution of all tasks set by ‘task’ directives
  - There are some restrictions in usage.



# Implicit Barriers

- Directives with implicit barriers
  - Parallel
    - Indispensible barriers
  - Do/for, Sections, Single, Workshare
    - Removable barriers
- Directives without implicit barriers
  - Master, Task

# Other Library Routines

- Rarely used but useful ones
  - `omp_get_wtick`
  - `omp_in_parallel`
  - `omp_set_dynamic`, `omp_get_dynamic`
  - `omp_set_nested`, `omp_get_nested`
  - `omp_set_schedule`
  - `omp_init_lock`, `omp_set_lock`, `omp_unset_lock`,  
`omp_destroy_lock`
  - Other lock routines

# References

- C. Evangelinos,  
Parallel Programming for Multicore Machines Using  
OpenMP and MPI
- R. Chandra, L. Dagum, D. Kohr, D. Maydan *et al.*,  
Parallel Programming in OpenMP
- B. Barney,  
OpenMP
- L. R. Scott, T. Clark, and B. Bagheri,  
Scientific Parallel Computing