

Parallel Scientific Computation

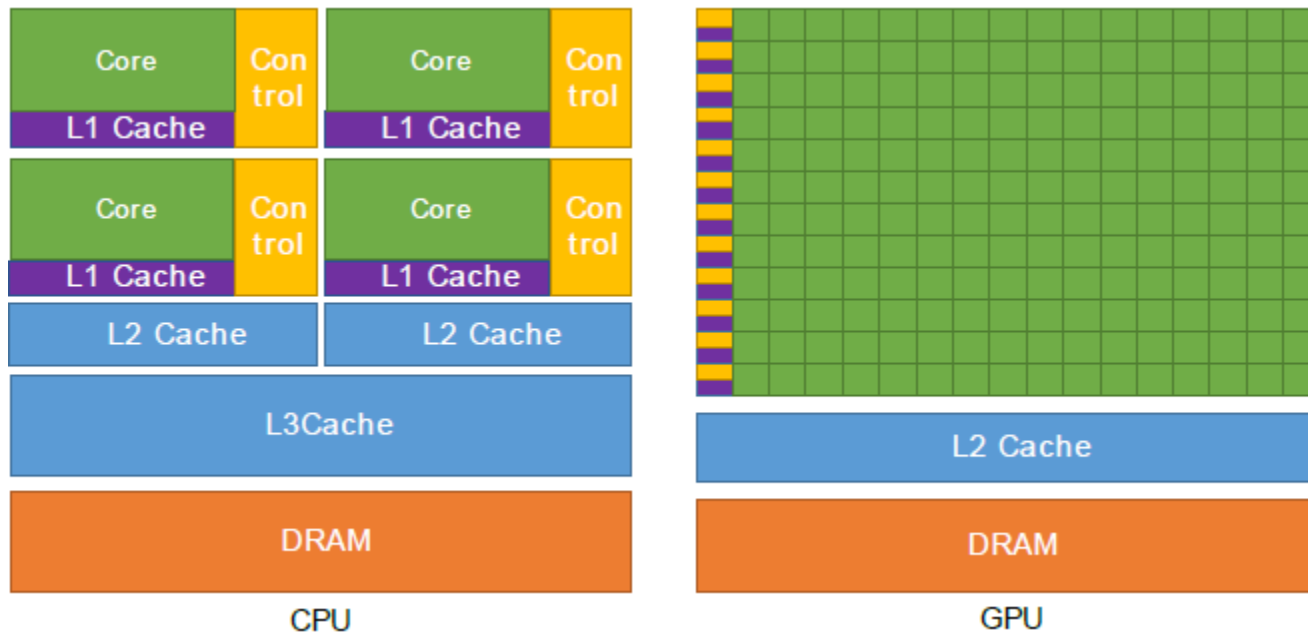
CUDA 1

J.-H. Parq

IPCST

Seoul National University

CPU vs GPU

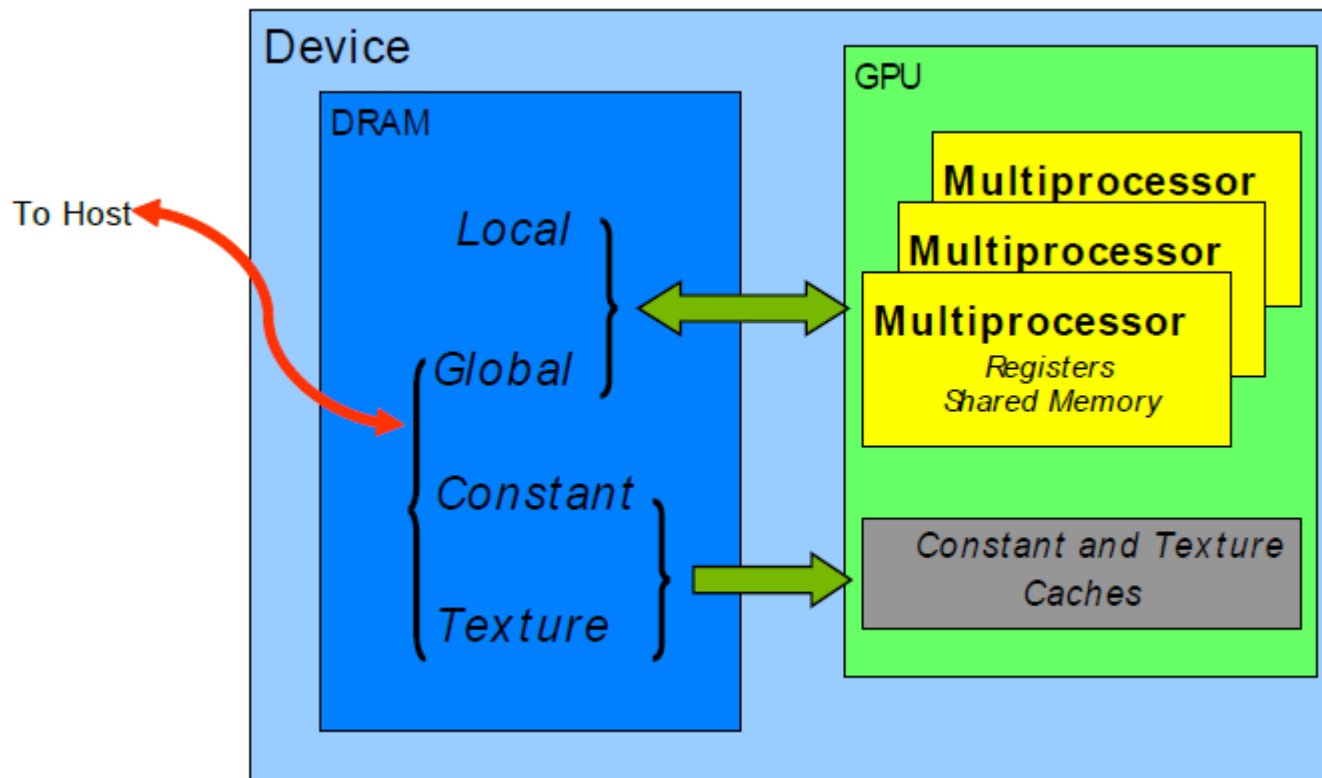


- GPU - many cores, small cache
 - Figure from NVIDIA CUDA C Programming Guide

What is GPGPU?

- General-Purpose computing on Graphics Processing Units
- Using GPUs not for graphics but for mathematical computation
- Transferring data to GPUs and computing
- Stream processing
 - Cf.) pipelining
- Interfaces: CUDA, OpenCL, OpenACC,

Memory Spaces on a GPGPU Device



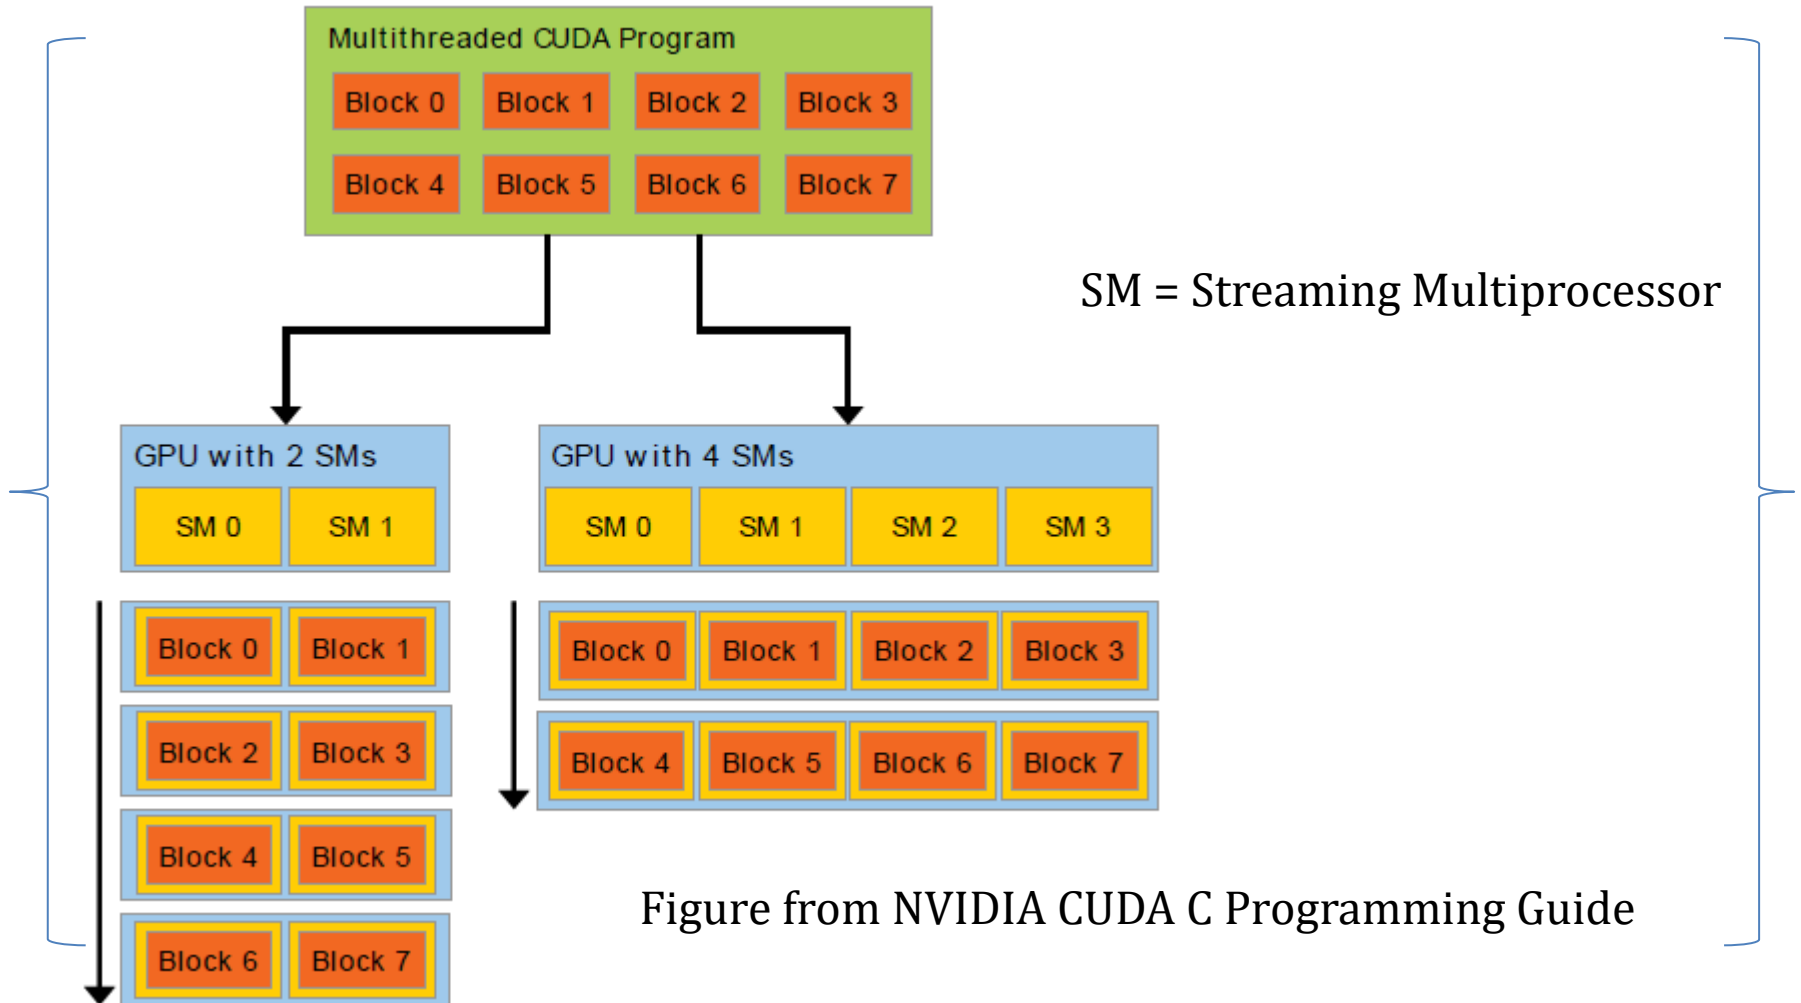
- Detailed structure depends on the GPU compute capability.

Figure from NVIDIA CUDA C Best Practices Guide

What is CUDA?

- Compute Unified Device Architecture
- General purpose parallel computing platform and programming model by NVIDIA
 - Only for Nvidia GPUs
- It provides a small set of extensions to C, C++, or Fortran.
 - Wrappers are available for Python, MATLAB, R, ...
- It supports heterogeneous computing using both CPUs and GPUs.
- SIMT (Single Instruction Multiple Threads)
 - Scalable programming model

Scalable Programming



CUDA Programs

- CUDA C
 - Extension: .cu
 - Compiling
 - nvcc xx.cu
 - See 'CUDA Compiler Driver NVCC' for options
- CUDA Fortran
 - Extension: .cuf, .CUF (+preprocessor)
 - Based on F90 or Fortran 2003
 - Compiling
 - NVIDIA HPC compiler: nvfortran xx.cuf
 - Portland group compiler: pgf90 xx.cuf / pgfortran xx.cuf

Data Transfer

- Between the host CPU and the device GPU
- In C,
 - `cudaMemcpy(dev_ptr, host_ptr, size, cudaMemcpyHostToDevice)`
 - `cudaMemcpy(host_ptr, dev_ptr, size, cudaMemcpyDeviceToHost)`
- In Fortran
 - `Assignment` statements
 - You can also use `cudaMemcpy` in the runtime routines.

Data Transfer Simple Example

CUDA Fortran

.....

Use **cudafor**

Integer , parameter :: N = 1024

Real :: H_a(N), H_b(N)

Real, **device** :: D_a(N)

.....

D_a = H_a

.....

H_b = D_a

.....

CUDA C

.....

size_t size = N * sizeof(float);

float *h_a = (float*)malloc(size);

float *d_a; **cudaMalloc(&d_a, size);**

.....

cudaMemcpy(d_a, h_a, size,
 cudaMemcpyHostToDevice);

.....

cudaMemcpy(h_a, d_a, size,
 cudaMemcpyDeviceToHost);

cudaFree(d_a);

.....

Kernels

- User-defined subroutines executing on GPUs

- CUDA Fortran

```
Attributes(global) subroutine k(a,b)
```

```
.....
```

```
End subroutine
```

```
Program .....
```

```
.....
```

```
Call k<<<M, N>>>( a, b )
```

```
.....
```

```
End Program .....
```

- CUDA C

```
__global__ void k(float *a, float *b){
```

```
.....
```

```
}
```

```
int main() {
```

```
.....
```

```
k<<<M, N>>>(a, b);
```

```
.....
```

```
}
```

Thread Hierarchy

- Threads in one block
 - An example in CUDA C Programming Guide

```
__global__ void VecAdd(float* A, float* B, float* C)
```

```
{  
    int i = threadIdx.x;  
    C[i] = A[i] + B[i];  
}
```

```
int i = blockIdx.x;
```

Fortran

Integer, value :: I

I = threadIdx%x

I = blockIdx%x

```
int main()
```

```
{  
    ...  
    VecAdd<<<1, N>>>(A, B, C);  
    ...  
}
```

Total threads in one block (= blockDim.x)

blockDim%x Fortran

```
VecAdd<<<M, 1>>>(A, B, C);
```

Fortran

Call VecAdd<<<1, M>>>(A, B, C)

Call VecAdd<<<M, 1>>>(A, B, C)

of blocks = 1

Thread Hierarchy

- Grid of thread blocks
 - 2D example (CUDA C)
 - threadIdx.y = 0 ~ 2
 - blockIdx.y = 0 ~ 1 *blockDim.y = 3*
 - 3D is also possible.

```
main()
{
    ...
    dim3 TPB(4,3);
    dim3 nBlocks(3,2);
    kernel<<<nBlocks, TPB>>>(A, B, C);
    ...
}
```

*Usually (16,16)
i.e., 256 threads
per block*

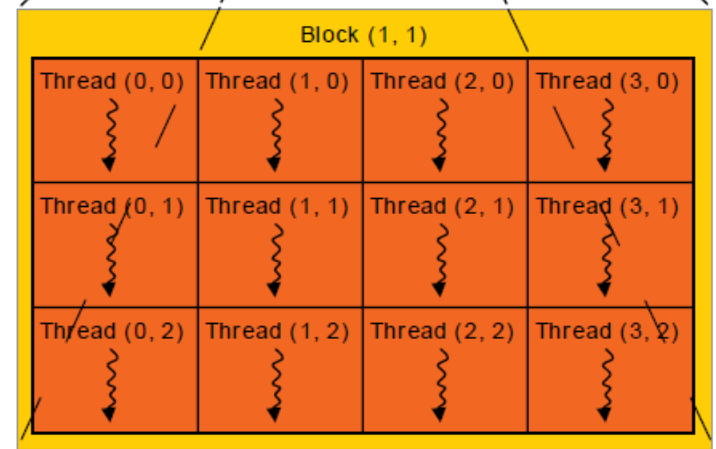



Figure from NVIDIA CUDA C Programming Guide

Thread Hierarchy

- Grid of thread blocks

- 2D example (CUF)

- $\text{threadIdx}\%y = 0 \sim 2$

- $\text{blockIdx}\%y = 0 \sim 1$

$\text{blockDim}\%y = 3$

...

USE CUDAFOR

...

`type (dim3) :: TPB, GRID`

`TPB = dim3(4,3,1)`

`GRID = dim3(3,2,1)`

`Call kernel<<<GRID, TPB>>>(A, B, C)`

...

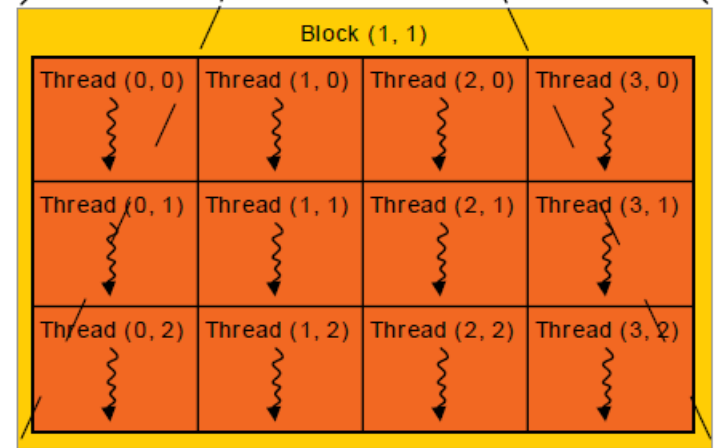



Figure from NVIDIA CUDA C Programming Guide

Memory Hierarchy

- Multiple memory spaces

- Speed

Global < Local < Shared < Regs.

- Capacity

Regs. < Shared < Local < Global

❖ Registers ~ hidden local memory

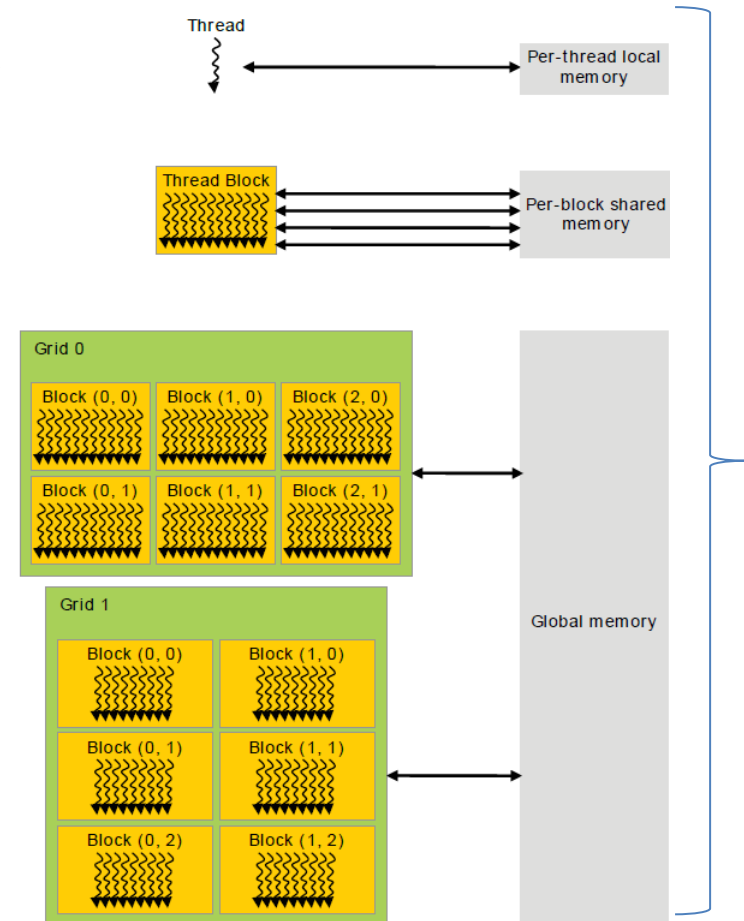


Figure from NVIDIA CUDA C Programming Guide

General CUDA Program Structure

1. Initialization

- Usually implicit in the program
- Selecting a GPU to run on

2. Global memory allocation on the GPU

- Maybe implicit or not

3. Data transfer: Host (usually CPU) → GPU

4. Launch GPU kernels from the host

5. Setting variables & GPU computation

- Variables: Local on a thread or Shared by a block

6. Result transfer: GPU → Host (usually CPU)

7. Output & memory freeing on the GPU

- Freeing may be implicit or not

Matrix Transpose (CUDA C)

```
#include <stdio.h>
#define nTx 16          // This should be 16 x n
#define nTy 16
main()
{
    .....
    cudaMalloc(&inmat_d, size);
    cudaMalloc(&outmat_d, size);
    nBx = Nrow/nTx + (Nrow%nTx != 0);
    nBy = Ncol/nTy + (Ncol%nTy != 0);
    dim3 grid(nBx, nBy), threads(nTx, nTy);
    cudaMemcpy(inmat_d, inmat_h, size, cudaMemcpyHostToDevice);
```


Matrix Transpose (CUDA C)

```
transpose<<<grid, threads>>>(inmat_d, outmat_d, Nrow, Ncol);  
cudaMemcpy(outmat_h, outmat_d, size, cudaMemcpyDeviceToHost);
```

```
.....  
}
```

```
__global__ void transpose(double *inmat_d, double *outmat_d, int Nrow,  
    int Ncol)
```

```
{  
    int i = blockIdx.x * blockDim.x + threadIdx.x;  
    int j = blockIdx.y * blockDim.y + threadIdx.y;
```

```
    outmat_d[j*Nrow + i] = inmat_d[i*Ncol + j];
```

```
}
```

Matrix Transpose (CUF)

```
program MT
  use cudafor
  use MTkernel
  .....
  type(dim3) :: grid, threads
  .....
  grid = dim3(ceiling(real(nrow)/ntx), ceiling(real(ncol)/nty), 1)
  threads = dim3(ntx,nty,1)
  inmat_d = inmat_h
  call transpose<<<grid, threads>>>(inmat_d, outmat_d, nrow, ncol)
  outmat_h = outmat_d
  .....
end program MT
```

Matrix Transpose (CUF)

Module MTkernel

attributes(global) subroutine transpose(inmat_d, outmat_d, nrow, &
ncol)

real(8), intent(in) :: inmat_d(nrow, ncol)

real(8), intent(out) :: outmat_d(ncol, nrow)

integer :: i, j

i = (blockIdx%x-1) * blockDim%x + threadIdx%x

j = (blockIdx%y-1) * blockDim%y + threadIdx%y

outmat_d(j, i) = inmat_d(i, j)

end subroutine transpose

End Module MTkernel

Coalesced Memory Access

- All threads in a block or of a warp should access global memory in a range, no same address, fitting in a cache line.
 - Warp: a group of 32 threads in SIMT architecture
 - The coalescing condition depends on the compute capability of the GPU.
 - Good access
 - Ex.) 4B words
- Coalesced access






Figure from NVIDIA CUDA C Best Practices Guide

Coalesced Memory Access

- A little bit bad access
 - Ex.) misaligned sequential addresses within five 32B segments



- Bad access



- Shared memories can increase access efficiency.

Figures from NVIDIA CUDA C Best Practices Guide

Matrix Transpose (CUDA C)

```
__global__ void transpose(double *inmat_d, double *outmat_d, int Nrow,  
    int Ncol)  
{  
    __shared__ double tile[nTx][nTy]  
    int bx = blockIdx.x * blockDim.x ;  
    int by = blockIdx.y * blockDim.y ;  
    int i = threadIdx.x;  
    int j = threadIdx.y;  
  
    tile[i][j] = inmat_d[(bx+i)*Ncol + by+j];  
    __syncthreads();  
    outmat_d[(by+i)*Nrow + bx+j] = tile[j][i];  
}
```

Matrix Transpose (CUF)

Module Mtkernel

Integer, parameter :: ntx = 16

Integer, parameter :: nty = 16

Contains

Attributes(global) Subroutine transpose(inmat_d, outmat_d, &
nrow, ncol)

Real(8), intent(in) :: inmat_d(nrow, ncol)

Real(8), intent(out) :: outmat_d(ncol, nrow)

Real(8), shared :: tile(ntx, nty)

Integer :: i, j

Matrix Transpose (CUF)

```
i = (blockIdx%x-1) * blockDim%x + threadIdx%x  
j = (blockIdx%y-1) * blockDim%y + threadIdx%y  
tile(threadIdx%x, threadIdx%y) = inmat_d(i, j)
```

Call `syncthreads()`

```
i = (blockIdx%y-1) * blockDim%y + threadIdx%x  
j = (blockIdx%x-1) * blockDim%x + threadIdx%y  
outmat_d(i, j) = tile(threadIdx%y, threadIdx%x)
```

End Subroutine transpose

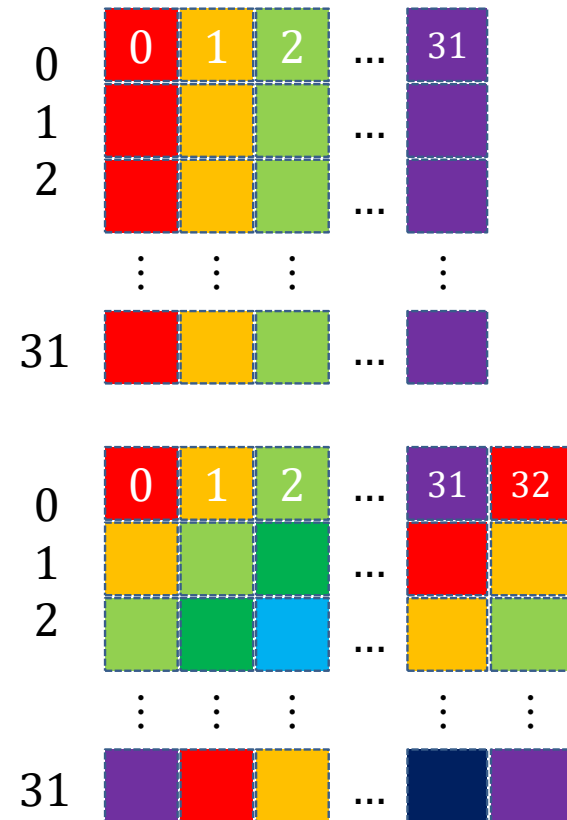
End Module MTkernel

Bank Conflicts

- **Bank:** equally-sized memory units in shared memory. Simultaneously accessed.
 - 32 banks for compute capability 3.x or higher
- Bank rules
 - 1 thread 1 bank access → OK
 - 2 thread 1 bank & the same word → OK
 - Broadcast reading
 - 2 thread 1 bank & different words → Conflict!
- ❖ Word: 32-bit mode or 64-bit mode

Bank Conflicts

- 2-D array example
 - 32 x 32 array (C)
 - Red: bank 0
 - Orange: bank 1
 - ...
- Remedy: array padding
 - `__shared__ double tile[32][33]`
 - `Real(8), shared :: tile(33, 32)`



Kernel Do Directive in CUF

- Similar to the OpenMP do directive
 - See more cases in the references.

- Ex. 1)

```
!$cuf kernel do(2) <<< (*,*), (32, 4)>>>
```

```
Do j = 1, m
```

```
  Do i = 1, n
```

```
    a(i,j) = b(i,j) + c(i,j)
```

```
  End do
```

```
End do
```

- (2): collapse
- (32,4): 32 thread ID x → inner loop
4 thread ID y → outer loop

- Ex. 2)

```
sum = 0.0
```

```
!$cuf kernel do <<<*,*>>>
```

```
Do i = 1, nx
```

```
  sum = sum + d_a(i)
```

```
End do
```

❖ Automatic reduction

➤ *,* : set by the compiler

※ Make sure to check restrictions.

OpenACC

- Acronym for **open accelerators**
- API to provide **easy** parallel programming for accelerators
 - **Similar to OpenMP**
 - Not only for Nvidia GPUs but also for AMD GPUs
- Supported languages: Fortran, C, C++
- Compilers
 - Commercial / semi-commercial (NVIDIA) / academic (from universities) / free (GCC)

OpenACC

- Directives
 - Similar to OpenMP
 - `#pragma acc` or `!$acc`
- Parallelizing loops
 - ‘parallel loop’ directive: similar to ‘parallel do/for’
 - ‘kernels’ directive: automatic optimization
 - Depending on your compiler
 - ‘loop’ directive: direction for the following loop
 - Used in ‘parallel’ or ‘kernels’ blocks
 - Clauses: `private`, `reduction`, `collapse`,
 - ‘kernels loop’ directive

OpenACC Examples

Fortran

!acc kernels

Do i=1, N

 x(i) = i

 y(i) = 0

End do

Do i=1, N

 y(i) = x(i)*3.0 + y(i)

End do

!acc end kernels

C

#pragma acc kernels

{

 for (i=0, i<N, i++) {

 x[i] = (float)i;

 y[i] = 0.0f;

 }

 for (i=0, i<N, i++)

 y[i] += 3.0f*x[i];

}

OpenACC Examples

Fortran

!\$acc parallel loop

do j=1,m-2

!\$acc loop

do i=1,n-2

A(i,j) = Anew(i,j)

end do

end do

C

#pragma acc parallel loop

for (int j = 1; j < n-1; j++)

{

#pragma acc loop

for (int i = 1; i < m-1; i++)

{

A[j][i] = Anew[j][i];

}

}

OpenACC

- **Implicit data transfer** at the beginning and the end of 'kernels' or 'parallel' blocks → **overhead**
- **Explicit data transfer**
 - '**data**' directive: defining a data block
 - Controlling memory allocation on a device at the block head & Data transfer depending on clauses
 - Clauses: **copy**(=copyin+copyout), **copyin**(host→device), **copyout**(host→device), **create**, **present**, **deviceptr**
 - **No data transfer** for the variables at the front and the end of 'kernels' or 'parallel' blocks **in the data block**
 - '**update**' directive: data synchronization (no block)
 - Clauses: **device**(host→device), **self**(device→host)

OpenACC

- Example of 'data' directive (C)

```
#pragma acc data copy(A), create(Anew)
```

```
{
```

```
while ( err > tol && iter < iter_max ) {
```

```
    err=0.0;
```

```
    #pragma acc parallel loop reduction(max:err)
```

```
    for( int j = 1; j < n-1; j++) {
```

```
        for(int i = 1; i < m-1; i++) {
```

```
            Anew[j][i] = 0.25 * (A[j][i+1] + A[j][i-1] + A[j-1][i] + A[j+1][i]);
```

```
            err = max(err, abs(Anew[j][i] - A[j][i]));
```

```
        }
```

```
    }
```

```
.....
```

OpenACC

- Another example of 'data' directive (Fortran)
 - Clause: p... = present or ...
 - Ex.: pcopy = present or copy

```
!acc data pcopyout(x(1:M)), pcreate(y(1:M))
```

```
!acc parallel loop
```

```
Do i = 1, N
```

```
  x(i) = i
```

```
  y(i) = 0
```

```
End do
```

```
.....
```

```
!acc parallel loop
```

```
Do i = 1, N
```

```
  y(i) = y(i) + x(i)
```

```
End do
```

```
!acc end data
```

References

- CUDA C Programming Guide
 - CUDA C++ Programming Guide
- CUDA C Best Practices Guide
 - CUDA C++ Best Practices Guide
- CUDA Fortran Programming Guide
- Wikipedia

References

- Introduction to CUDA Fortran
 - GPU Technology Conference 2013
- OpenACC homepage
<https://www.openacc.org/>
- OpenACC getting started guide
- J. Sanders & E. Kandrot, CUDA by Example